



Haribote File System Specification

Version 1.0

Written by Allen He

Notices:

This document may contain errors, and the author bears no full responsibility for these errors.

No rights are granted to create modifications or derivatives of this document. But you may publish an unmodified copy.

The author reserves the right of final interpretation of this document.

© Allen He.

Thanks Hidemi KAWAI, for guiding so many beginners in writing operating system.

Contents

| | | |
|----------|--------------------------------|-----------|
| 1 | Introduction | 1 |
| 2 | Structure | 2 |
| 2.1 | Overall Structure | 2 |
| 2.2 | Boot Sector | 3 |
| 2.3 | Super Sector | 4 |
| 2.4 | Segment Super Sector | 5 |
| 2.5 | Inode | 6 |
| 2.6 | Directory | 7 |
| 2.7 | Additional Block | 7 |
| 3 | Operating Specification | 8 |
| 3.1 | Format | 8 |
| 3.2 | Create a file | 9 |
| 3.3 | Create a directory | 10 |
| 3.4 | Delete a file | 10 |
| 3.5 | Delete a directory | 11 |
| 3.6 | Read a file | 11 |
| 3.7 | Write to a file | 11 |
| 3.8 | Notes | 12 |
| 4 | Parameter Specification | 13 |

1

Introduction

Haribote File System (HAFS) is a lightweight file system designed for the Haribote OS and Haribote OS for Metaverse.

HAFS inherits the minimalist style of Haribote OS, but mature file systems contain complex structures. Therefore, to enhance robustness and utility, HAFS must incorporate certain moderately complex structures.

HAFS consists of a Boot Sector, a Super Sector, bitmaps and segments. The smallest unit of storage allocation is the block. And each segment contains a fixed number of contiguous blocks. Every segment consists of Segment Super Sector, segment block maps, segment inode maps, inodes and data blocks.

Like Unix-like operating systems' file system, HAFS uses inodes to record file information. But HAFS's inodes are smaller, and other attributes are recorded in additional blocks.

This document provides a detailed technical specification of HAFS, covering its structure, implementation, and parameters.

2

Structure

2.1 Overall Structure

| | | | | | | |
|-------------|------------------|-------------|-----------|-----------|-----------|-----|
| Boot Sector | Reserved sectors | Super block | Block map | Inode map | Segment 0 | ... |
|-------------|------------------|-------------|-----------|-----------|-----------|-----|

Table 1: The structure of HAFS.

Table 1 shows the structure of HAFS. The partition is divided into segments. There are five special parts in the partition: Boot Sector, Reserved sectors, Super block, Global block bitmap and Global inode map. Despite being global in scope, these structures are physically stored within segment 0.

The global block bitmap identifies segments with free blocks; likewise, the global inode bitmap identifies segments with free inodes.

| | | | | |
|----------------------|-----------|-----------|--------|-------------|
| Segment Super Sector | Block map | Inode map | Inodes | Data blocks |
|----------------------|-----------|-----------|--------|-------------|

Table 2: The structure of a segment.

Each segment begins with Segment Super Sector, next are block map, inode map and inodes. The remaining part is the data blocks.

Unlike other segments, the Segment Super Sector of segment 0 resides following the global inode bitmap.

2.2 Boot Sector

Following is the structure of Super Sector:

| Offset | Name | Length | Describe |
|--------|-----------------------|--------|---------------------------------------|
| 0x0 | Jump instruction | 3 | Jump instruction |
| 0x3 | OEM code | 8 | Set to "HAFS " |
| 0xB | Sector size | 2 | Sector size in bytes |
| 0xD | Block size | 1 | Block size in sectors |
| 0xE | Reserved sector | 2 | Reserved sector number |
| 0x10 | Unused | 4 | Unused |
| 0x14 | Media descriptor byte | 2 | Disk type, 0xF8 for Harddisk |
| 0x16 | Unused | 4 | Unused |
| 0x1A | Signature | 4 | The signature, set to 0x00800080 |
| 0x1E | Sector number | 8 | Sector number in total |
| 0x26 | Version | 4 | Version number, set to 0x1 |
| 0x2A | Super Sector start | 4 | The start sector* of the Super Sector |
| 0x2E | Boot code | — | Boot code |
| 0x1FE | Boot sector signature | 2 | 0xAA55 |

*: The starting sector is relative to the partition.

2.3 Super Sector

Following is the structure of Super Sector:

| Offset | Name | Length | Describe |
|--------|-------------------|--------|--------------------------|
| 0x0 | Signature | 4 | ”HAFS” |
| 0x4 | Partition size | 8 | Partition size in byte |
| 0xC | Block size | 2 | Block size in byte |
| 0xE | Block map start | 2 | Block map start sector* |
| 0x10 | Inode map start | 2 | Inode map start sector* |
| 0x14 | Segment 0 start | 4 | Segment 0 start sector* |
| 0x18 | Root inode | 4 | Root directory's inode |
| 0x1C | Disk type | 4 | Disk type |
| 0x1C | 0x1 | | Hard disk |
| 0x1C | 0x2 | | CD / DVD |
| 0x1C | 0x4 | | Floppy |
| 0x1C | 0x8 | | USB Flash Disk |
| 0x1C | 0x10 | | Bootable |
| 0x1C | 0x20 | | Hidden |
| 0x1C | 0x40 | | Read-only |
| 0x20 | Inode per segment | 4 | Inode number per segment |
| 0x24 | Segment number | 4 | Segment number in total |
| 0x28 | Block map size | 4 | Block map size in byte |
| 0x2C | Inode map size | 4 | Inode map size in byte |
| 0x30 | Check code | 4 | CRC32 check code |

*: The starting sector is relative to the partition.

2.4 Segment Super Sector

Following is the structure of Segment Super Sector:

| Offset | Name | Length | Describe |
|--------|-------------------|--------|----------------------------------|
| 0x0 | Signature | 4 | "HAFS" |
| 0x4 | Block number | 2 | The block number of this segment |
| 0x6 | Block map start | 2 | Block map start sector* |
| 0x8 | Inode map start | 2 | Inode map start sector* |
| 0xA | Inode area start | 2 | Inode area start sector* |
| 0xC | Free block number | 4 | Free block number |
| 0x10 | Free inode number | 4 | Free inode number |
| 0x14 | Check code | 4 | CRC32 check code |

*: The starting sector is relative to the **segment**.

2.5 Inode

Following is the structure of inode:

| Offset | Name | Length | Describe |
|--------|----------------------------|--------|---------------------------|
| 0x0 | Type | 4 | The type of this inode |
| 0x1 | | | Present |
| 0x2 | | | File |
| 0x4 | | | Directory |
| 0x8 | | | Hidden |
| 0x10 | | | Read-only |
| 0x20 | | | System |
| 0x40 | | | With long filename |
| 0x80 | | | With additional attribute |
| 0x4 | Size | 8 | File size in byte |
| 0xC | Additional attribute block | 4 | * |
| 0x10 | Direct blocks | 12*4 | * |
| 0x40 | L1 indirect blocks | 8*4 | * |
| 0x60 | L2 indirect blocks | 4*4 | * |
| 0x70 | L3 indirect blocks | 4*4 | * |

*: The starting block is relative to the partition.

2.6 Directory

Following is the structure of directory:

| Offset | Name | Length | Describe |
|---------------|-------------|---------------|----------------------------|
| 0x0 | Next record | 4 | The offset of next record* |
| 0x4 | Inode | 4 | The inode of the file |
| 0x8 | Name | — | File name |

*: The offset is relative to this record.

2.7 Additional Block

Following is the types of additional attribute:

| Number | Name | Describe |
|---------------|--------------|-------------------------------|
| 0x2 | Long name | The long name of file |
| 0x8 | Father inode | The inode of parent directory |
| 0x10 | Create time | Create time |
| 0x20 | Access time | Access time, only for files |
| 0x40 | Modify time | Modify time, only for files |

3

Operating Specification

3.1 Format

To format a partition to HAFS, follow the steps below:

1. Collect and calculate the necessary parameters, including segment number, the starting sector numbers of maps and the starting sector number of the zero segment.

(After completing each of the following steps, write the changes into the partition immediately.)

2. Construct the Boot Sector.
3. Construct the Super Sector, fill the area preceding the check code, and then calculate the check code for the filled content.
4. Set the global block map and the global inode map to zero.
5. Construct the Segment Super Sector of the zero segment, fill the area preceding the check code, and then calculate the check code for the filled content.
6. Set the block map and the inode map for the zero segment; the bits corresponding to the blocks used (i.e. the blocks before the zero segment, the Segment Super Sector and maps) should be set.
7. Construct the root inode. Set the inode type "Present" and "Directory".
8. Construct the Segment Super Sectors and maps of all other segments.

3.2 Create a file

To create a file, follow the steps below:

1. Check the validity of the destination path.
2. Get the inode number of the destination directory.
3. Traverse the directory entries to check for any duplicate files.
4. Allocate a new inode and add a new record to the destination directory.
4. Initialize the inode and fill the create time, the access time and the modify time.

3.3 Create a directory

To create a directory, follow the steps of creating a file. Notice that directories don't have the access time and the modify time.

3.4 Delete a file

To delete a file, follow the steps below:

1. Check if the file exists.
2. Get the inode number of the file and the file's parent directory.
3. Remove the record of the file in the parent directory.
4. Free all blocks allocated by the file.
5. Set the inode of the file to zero.
6. Free the inode of the file.

3.5 Delete a directory

To delete a directory, follow the steps of deleting a file. Notice that directories should be empty.

3.6 Read a file

To read a file, check if the file exists first, then read the data of the file, and change the access time of the file to current time at last.

3.7 Write to a file

To write a file, follow the steps below:

1. Check if the file exists.
2. Calculate the number of the free blocks, and then check if it is enough to allocate blocks for the new data.
3. Change the access time and the modify time of the file to current time.
4. Allocate new blocks for the file.
5. Put the data to the blocks.

3.8 Notes

In all operations except the format, write the changes to the buffer of the operating system if it exists, or write them into disk immediately.

The number of free blocks/inodes should be checked before allocation.

The access and modify time should be changed in advance to prevent accidental loss.

4

Parameter Specification

Here are some fixed parameters of HAFS:

1. Block size: 1024/2048/4096 bytes
2. Segment size: 1024 blocks
3. Maximum partition size: LLONG_MAX KiB

Here are some alignment requirements:

The Super Sector and the Segment Super Sectors: Align to sector.

Maps: Align to sector.

Zero segment start: Align to block.

Data part: Align to block.